



Edge Deployment: Plant Health Classification with MobileNet V3 on a Raspberry Pi Zero 2 W


Yusuf Savaş | Advisor: Fikret Gürgen
Boğaziçi University | SWE 599 Graduate Project

The Problem

Pests and diseases destroy up to **40% of global crops** every year, about **US \$220 billion** in losses. Early, on-site detection from leaf images, protects yield and cuts blanket pesticide use.


Goal: a self-contained classifier that labels a live camera frame as **Healthy**, **Diseased**, or **Background** — running entirely on a \$15, 512 MB Raspberry Pi Zero 2 W with *no* cloud, GPU, or NPU.

System Architecture



Camera (or stored images) → on-device preprocess → ONNX Runtime → web portal streamed over the LAN.

Dataset — 3 classes, 12,606 images




PlantVillage leaves collapsed to healthy / diseased; **COCO 2017** + field negatives form the background class that lets the model reject non-leaf scenes. Split **70 / 15 / 15**.


Key Design Choices

- 3-class topology** — a built-in *background* class rejects empty scenes with no extra detection model.
- FP32 over INT8** — the Cortex-A53 lacks *SDOT/UDOT*, so INT8 falls back to slow software; FP32 uses mature NEON FMA instructions.
- Native C++ runtime** — keeps memory far under the 512 MB ceiling; zero-allocation hot loop.
- MobileNet V3-Small** — only **1.52 M** parameters.

Training Pipeline



On-Device Runtime (C++)




Edge Benchmarks

Build	Split	Acc.	img/s
FP32 (4 thr.)	3-class	99.81%	15.73
FP32 (1 thr.)	2-class	99.83%	8.03
INT8 (4 thr.)	3-class	97.05%	9.83
INT8 (1 thr.)	3-class	97.21%	6.04

Per-class F1: Healthy 0.995, Diseased 0.998, Background 0.999. INT8 is slower here — wrong tool for this silicon.

Live Demo & Test Bench



Web portal shows live prediction, per-class confidence, FPS, and CPU temperature. A 3D-printed PLA bench fixes camera height and lighting for repeatable physical validation.

Conclusion

A native C++ end to end pipeline runs MobileNet V3 at **~8 img/s** on a Pi Zero 2 W (live demo: **inference latency ~68 ms**, **encoding latency ~31 ms** per frame). The background class removes the need for a second model, and FP32 beats INT8 on this CPU. *Future work:* fine-tune on deployment-camera captures and add field validation in live greenhouses.

Headline Results

68	31	8
ms inference (live)	ms encoding (live)	images / sec on the Pi